

RAZEEN H

GENERATIVE AI ENGINEER

razeen.official23@gmail.com | **INR :** +91 9345035514 | **UAE :** +971 50 769 7208

linkedin.com/in/razeen-developer | github.com/Razeen23 | razeen-portfolio.vercel.app

Professional Summary

Full Stack Developer with 3+ years shipping Generative AI features on a production hosting platform. Support answers were not tied to the internal docs, so I built a retrieval-augmented assistant with LangChain and FAISS. OpenAI and Google Gemini answer only from that retrieved text. The assistant also connects to the server to install packages, create files and folders, and run DNS, SSL, and stack checks, and it opens a ticket only when it cannot finish the fix. Those features sit on the PHP, JavaScript, and MySQL product used by 5,000+ users.

Skills

Generative AI: Retrieval-Augmented Generation (RAG), Large Language Models, LangChain, FAISS, embeddings, vector search, prompt engineering

Models and APIs: Ollama, OpenAI API, Google Gemini API, FastAPI, Python

Application stack: PHP, JavaScript, AJAX, React.js, TypeScript, Streamlit, Node.js, Express.js

Data: MySQL, PostgreSQL, SQLite, Pandas

Delivery: Docker, Kubernetes, Helm, Harbor, GitHub Actions, Git, Linux

Work Experience

Full Stack Developer, Hostzop Cloud Technologies Pvt. Ltd. Aug 2023 - Present

- Problem: hosting and server questions were answered from memory, not from the internal docs. I built a retrieval-augmented assistant with LangChain and FAISS, and FastAPI for retrieval and the reply. OpenAI and Google Gemini answer only from the retrieved text. Impact: the reply stays inside the docs.
- Problem: installing a package, creating a folder, and DNS, SSL, and stack checks were done by support by hand. I connected the assistant to the server so it can do those actions itself. Impact: a ticket is opened only when the assistant cannot finish the fix.
- Problem: one question about an earnings-call transcript meant reading the whole document, and the summary was written by hand. I built document ingestion, vector search, and a Streamlit page for questions and summaries. Impact: the answer and the summary come from the transcript.
- Problem: each lead in a spreadsheet needed a personal email, and the same person could be emailed twice. I drafted the mail with a local Ollama model, logged each send in SQLite, and sent it over SMTP with a rate limit. Impact: no lead is emailed twice.
- The assistant and the product around it run on the PHP, JavaScript, and MySQL platform used by 5,000+ users, released with Docker, Harbor, Helm, Kubernetes, and GitHub Actions.

Full Stack Developer Intern, Aspire Systems Apr 2023 - Jul 2023

- Engineered a real-time face recognition attendance system using Python, OpenCV, cvzone, and NumPy.
- Serialized a facial-matching model with Pickle and sent recognized attendance records to Firebase.

Projects

Generative AI Support Assistant

Python, FastAPI, LangChain, FAISS, OpenAI API, Google Gemini API

- Problem: a server or hosting question opened a ticket first, and the answer was not taken from the internal docs.
- Built: the assistant inside the hosting account. LangChain and FAISS pull the matching docs and runbooks. OpenAI and Google Gemini answer only from that text. FastAPI serves the reply.
- Built: the assistant connects to the server and carries out the fix. It installs packages, creates files and folders, and runs DNS, SSL, and stack checks.
- Impact: those steps no longer wait on a person. A ticket is opened only when the assistant cannot finish the fix.

AI Financial Document Analyzer

Python, Streamlit, vector search

- Problem: a question about an earnings-call transcript meant reading the whole document, and the summary was written by hand.
- Built: document ingestion and semantic retrieval over the transcripts, so the model reads the matching passages instead of the full file.
- Built: a Streamlit page where a user asks a question or requests a summary, and the reply is written from those retrieved passages.
- Impact: the answer and the summary come from the transcript, instead of a full reread and a hand-written summary.

AI Email Sender

Python, Ollama, SMTP, SQLite, Pandas

- Problem: each lead in a spreadsheet needed a personal email, and nothing stopped the same person from being emailed twice.
- Built: a local Ollama model drafts the email from the lead row. SQLite logs each send, and SMTP sends the mail with a rate limit.
- Impact: no lead is emailed twice, and the draft no longer depends on someone reading the sheet and writing the mail.

Production Platform for These AI Features

PHP, JavaScript, AJAX, MySQL, Docker, Kubernetes

- Problem: the assistant was useful only if it lived in the product customers already use, on the same release path as the rest of the platform.
- Built: the white-label portals, billing, and Stripe and Razorpay payments, and a Git-to-live path with Docker, Harbor, Helm, TLS, and DNS.
- Impact: the AI features run on the platform used by 5,000+ users, not as a separate demo.

SaaS Billing and Usage Dashboard

PHP, JavaScript, AJAX, MySQL, D3.js

- Problem: billing, usage, and server status sat in different places, and a paid order was turned on by hand.
- Built: the dashboards in PHP, JavaScript, AJAX, and D3.js. Billing, usage, and server metrics load in the page without a full reload. A reseller can attach their own domain and brand and sell hosting under their name.
- Built: Stripe and Razorpay confirm the payment inside the same dashboard and activate the service from that payment.
- Impact: the dashboards serve 5,000+ users and the white-label portals serve 1,000+ users. A paid order turns on with no manual step.

Education

Bachelor of Computer Applications (BCA) 2019 - 2022

Mohamed Sathak College of Arts and Science, Chennai

Diploma in Counseling Psychology (Online) 2020 - 2021

Sypa IIE

Certifications

- Full Stack Development - Offline Training
- Back-end Development - Meta
- JavaScript - Udemy
- Fundamentals of Cybersecurity - upGrad
- UX Design Process: Empathize, Define, and Ideate - Google

Awards

- Best Performer Award during the Aspire Systems internship for technical contribution and reliability.
- Led a 10-member team during NCC Camp, with responsibility for coordination and accountability.

Languages

English (Fluent), Tamil (Native), Hindi